

Complexity Signature of Generated Text

Ross Schmidt^{○a} and Ishanu Chattopadhyay^{○a}

This manuscript was compiled on February 4, 2026

We introduce a model-agnostic approach for detecting LLM-generated text by estimating the entropy rate of text mapped to symbol streams over a (26+space)=27-letter alphabet, and show that long-form outputs from LLMs exhibit systematically lower entropy rates than human-authored prose. Our Nonparametric Entropy-Rate Oracle (NERO) exploits this signal to achieve competitive, training-free separation of human and AI text, without requiring model access. When substring-frequency cutoff profile for the estimator is used as input to a supervised Gaussian process classifier, NERO achieves near-perfect detection accuracy (AUC = 98.9%), providing a principled, complexity-theoretic framework for ranking generative capacity and tracking changes in LLM behavior over time.

entropy rate | AI-generated text detection | large language models | algorithmic complexity | probabilistic automata

With the rise of generative artificial intelligence (AI), particularly large language models (LLMs), reliably distinguishing human-authored prose from machine-generated text is no longer trivial (1–3). Here we test the hypothesis that long-form human- and machine-generated text differ in statistical complexity when mapped to symbol streams over a 27-character alphabet (26 letters plus space) and treated as sample paths of a stochastic process (4), whose entropy rate quantifies intrinsic complexity. Most existing detectors (5, 6) instead rely on language-model scoring (e.g., likelihood/rank statistics), supervised classifiers over stylistic representations, or compressibility proxies, and therefore depend on reference model access, training, or calibration to evolving generators.

We introduce a nonparametric, learning-free entropy-rate estimator that operates directly on text, without model access, supervision, or retraining. Used as a threshold score, the resulting Nonparametric Entropy-Rate Oracle (NERO) achieves competitive, training-free discrimination, and reveals systematically lower entropy rates in contemporary LLM outputs than in human prose under a shared symbolization. Optionally, using the estimator's internal substring-frequency cutoff profile as features in a Gaussian process classifier yields near-perfect performance (AUC = 98.9%), surpassing model-driven baselines. Thus, NERO-estimated entropy rate functions as a model-agnostic statistic, akin to a physical quantity, for effective generative capacity, and provides a robust mechanism for AI text detection and a principled framework for ranking and tracking generative models over time. Since we target an intrinsic property of long-form text, task-based LLM benchmarks are not directly commensurate here (see Supplementary Information).

Results

To connect generative capacity with algorithmic complexity and then to entropy-rate of outputs, we begin by recalling the foundational concept of optimal two-part codes in algorithmic information theory (7). The Kolmogorov complexity $K(x)$ of a string x is defined as the length of the shortest program (in a fixed universal programming language) that produces x and halts. It provides a rigorous, machine-independent measure of the information content or compressibility of a string. Instead of describing x directly, we can choose to first describe a finite set S that contains x , and then identify x within S by its index in a standard enumeration of all items in S . This leads to the two-part code representation:

$$K(x) \leq K(S) + \log |S| + O(1), \quad [1]$$

where $K(S)$ is the complexity of describing the set S , and $\log |S|$ is the number of bits required to identify x within S . When x is a *typical* element of S , i.e., that it does not admit any significantly shorter description than most elements of S , the inequality becomes tight up to additive constants. The optimal two-part code is obtained by minimizing the sum $K(S) + \log |S|$ over all such sets containing x , yielding a minimal sufficient statistic for the data.

Author affiliations: ^aUniversity of Kentucky

RS carried out experimental runs and wrote the paper, IC conceived of research, implemented the algorithm, wrote the paper and procured support.

Authors have no competing interests.

²To whom correspondence should be addressed. E-mail: ishanu.ch@uky.edu

125 **Table 1. Cohorts and average entropy rates**

data source	mean H	median H	std. dev. H	count
GPT-3.5 (webform)	0.54	0.53	0.11	100
GPT-4o (webform)	0.62	0.63	0.09	98
GPT-4o (API)	0.64	0.66	0.09	136
Gemini (API)	0.66	0.67	0.10	987
Claude (API)	0.71	0.72	0.06	944
GPT-4.0 (webform+API)	0.71	0.71	0.09	82
GPT-5 (API)	0.74	0.74	0.08	197
Gutenberg project	0.77	0.78	0.12	4341
Arxiv papers	1.32	1.19	0.55	42

Now lets apply these notions to quantify the generative ability of AI agents (and humans). With each agent G_i we associate a set S_i that defines the collection of strings that the agent can possibly generate. Once S_i is fixed, we assume that the agent produces an individual output $x \in S_i$ using a standard sampling procedure. This shared decoding mechanism implies that the conditional complexity $K(x | S_i)$ is approximately constant across agents for typical strings. An equivalent way to view this assumption, is that agents draw from the same or similar set of possible strings, *i.e.*, if one agent can generate a particular string x , then so can the other, perhaps with different odds. Under either interpretation, the implication is that the conditional term $K(x | S_i)$ is effectively constant across agents, and since for typical x this term satisfies $K(x | S_i) \approx \log |S_i|$ (7), we have:

$$\log |S_1| = \log |S_2| + O(1). \quad [2]$$

Then, we have the proposition (See Methods for proof):

Proposition 1 (Two-Part Code to Entropy). *Let G_1 and G_2 be two generative processes outputting strings over a finite alphabet $x \in \mathcal{A}^n$, each respectively associated with a set S_1 and S_2 of possible outputs. Assuming $\log |S_1| = \log |S_2| + O(1)$, the following equivalence holds:*

$$K(S_1) > K(S_2) + O(1) \Leftrightarrow E_{x \sim G_1} [K(x)] > E_{x \sim G_2} [K(x)] + O(1).$$

Additionally, if G_i are stationary, then

$$K(S_1) > K(S_2) + O(1) \implies H(G_1) > H(G_2) + o(1), \quad [3]$$

where $H(G_i)$ is the Shannon entropy rate of generator G_i .

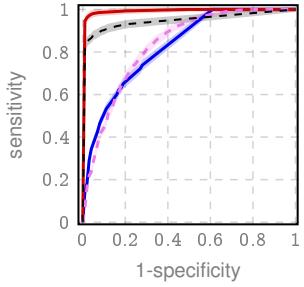
Proposition 1 motivates entropy rate as a discriminative statistic, suggesting recognition of AI-generated text by its systematically lower entropy rate. Estimating entropy rate for symbol streams is nontrivial (8). Accordingly, for NERO we adopt our prior nonparametric entropy-rate estimators based on probabilistic finite-state automata (PFSA) (9, 10) (See Supplementary Methods). With the standard approximate stationarity assumption (4) for long-form texts, Proposition 1 provides the key foundation of our claim.

Experiments

To evaluate our claim, we use a corpus (See Table 1), comprising human-authored texts (Project Gutenberg and arXiv) and long-form outputs from contemporary LLMs (GPT-3.5, GPT-4o, GPT-4.0, Claude, Gemini, and GPT-5), using a mix of API and web-form access (See Supplementary Methods). All documents were lowercased and mapped to a 27-symbol English-plus-space alphabet, removing punctuation, digits, and non-ASCII characters. For each document we

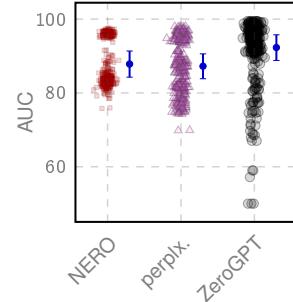
a. ROC (with 95% CI)

- NERO (no training): 82.7%
- NERO (with training): 98.9%
- - perplexity (HowGPT): 82.3%
- - ZeroGPT: 94.9%

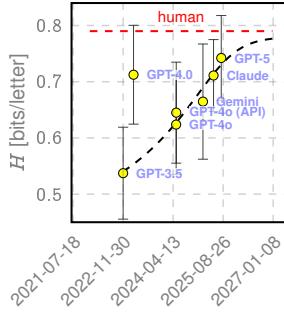


b. AUC (Cohort-Variation*)

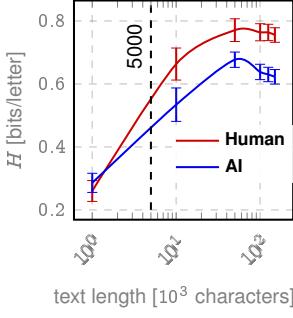
- *381 cohort combinations
- NERO-notraining
- △ perplexity (HowGPT)
- ZeroGPT



c. Birthtime Vs Performance



d. Length dependence



e. Topic dependence in human prose (Gutenberg)

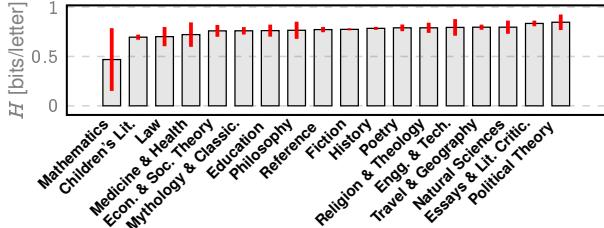


Fig. 1. NERO performance. a, Pooled human-AI ROC curves for training-free NERO (median entropy-rate estimate \hat{H}), and with a Gaussian-process classifier alongside baselines (perplexity from HowGPT and ZEROGPT scoring, See Supplementary Methods). b, Cross-cohort robustness across 381 out-of-sample discrimination tasks obtained by varying the human reference set (Gutenberg, arXiv, or both) and the subset of LLM generators. c, Mean entropy rates plotted at model release dates (“birthtime”), illustrating upward drift toward the human regime; a weighted Richards fit with a fixed ceiling at the mean human rate is overlaid for descriptive smoothing. d, Entropy-rate estimates versus document length under truncation (using Gutenberg texts for humans), indicating a practical lower bound for reliable discrimination at shorter lengths around 5,000 characters. e, Genre-stratified NERO entropy rates for human prose; error bars denote 95% intervals.

compute a family of entropy-rate estimates across substring-frequency thresholds $\{m_1, \dots, m_M\}$, where m specifies the minimum number of occurrences required for a substring to be retained in the PFSA construction (see Methods). In the training-free setting, we use the median across thresholded m -estimates (H). In the trained-NERO setting, the full m -dependence profile is used as features for training a Gaussian process classifier (11); importantly, the underlying entropy-rate estimator is unchanged. Using NERO, our estimates for arxiv papers (≈ 1.3 bits/letter, Table 1) closely align with classic estimates of English entropy under a 27-symbol alphabet (4), whereas Gutenberg prose exhibits lower

character-level entropy, consistent with greater redundancy in narrative text. Our experimental results (Fig. 1) demonstrate:

(1) *Training-free detection.* Using H directly as a detection score yields strong LLM-human discrimination without any training. (AUC = 0.824 on the pooled evaluation, Fig. 1a).

(2) *Training using substring-frequency threshold m-dependence.* We trained a Gaussian-process classifier on the m -dependent feature vector (see Methods) using a 50/50 train-test split, improving out-of-sample discrimination to AUC = 0.989 (Fig. 1a). For comparison, the perplexity baseline (HowGPT) attains AUC = 0.830, while ZEROGPT attains AUC = 0.946 on the same pooled evaluation.

(3) *Cross-cohort robustness.* Fig. 1b reports out-of-sample AUCs for $(2^2 - 1) \times (2^7 - 1) = 381$ distinct human-AI discrimination task obtained by varying the human reference and the subset of LLM generators included (See SI Methods). Across this combinatorial suite, training-free NERO maintains consistently strong performance, whereas trained baseline detectors exhibit substantially greater sensitivity to corpus composition and generator selection. This contrast underscores the NERO estimate as a calibration-free, model-agnostic statistic rather than a generator-dependent detector.

(4) *Temporal trajectory toward the human regime.* Fig. 1c plots cohort-level mean NERO-computed entropy rates at model release times with a weighted Richards (generalized logistic) fit and a fixed asymptotic ceiling given by the mean human rate (SI Methods), showing successive model generations shifting upward toward the human regime with diminishing returns. While not a calendar-date forecast, this supports the qualitative conclusion that model improvements are narrowing the entropy-rate gap. GPT-4.0 is a notable outlier, consistent with a regime shift induced by the architectural jump introduced in this model (multi-modal access and other experimental features) (12).

(5) *Length dependence.* Human-AI separation is robust for articles with > 10,000 characters (about the length of a full-length magazine article), with loss of discrimination possibly at lengths under 5,000 characters (Fig. 1d).

(6) *Genre dependence.* Formulaic human genres (Mathematical text, Children's Literature) attain lower entropy rates than expository genres (e.g. Literary criticism and political theory), defining the primary hard cases (Fig. 1e).

Discussion

We show that, under shared symbolization, long-form human-authored text exhibits a consistently higher estimated entropy rate than contemporary LLMs. In the two-part code perspective, when outputs are compared within a common representation, differences in entropy rate are consistent with differences in the effective descriptive complexity of the underlying generative mechanisms, suggesting that humans might operate with a richer, more complex internal model than current AI systems.

A complementary intuition follows from *Levin's universal semimeasure* $m(x)$ (13), which lower-bounds the probability assignable to strings by any computable process. It defines a universal prior by summing over all prefix-free programs p (14) that produce x on a universal Turing machine U (15):

$$m(x) = \sum_{p: U(p)=x} 2^{-|p|}. \quad [4]$$

and Levin's Coding Theorem links to Kolmogorov complexity,

$$K(x) = -\log m(x) + O(1), \quad [5]$$

which implies that strings with low $m(x)$ are algorithmically complex and lie in the deep tails of any computable generative process. Humans operating with an internal generative model S_{human} that likely possesses significantly higher descriptive complexity than models used in AI systems, are able to generate outputs that lie deeper in the algorithmic tail of $m(x)$, i.e., outputs that are rarer, more contextually novel, and of higher complexity. In contrast, LLMs are trained on finite, empirical datasets where such rare strings are underrepresented (by definition of being rare); since strings with low $m(x)$ are, by definition, uncommon, they are statistically unlikely to appear with sufficient frequency in training corpora. This possibly limits the effective model complexity $K(S_{\text{AI}})$, and hence constrains the AI's ability to reproduce or extrapolate into these low-probability regions.

Thus, despite the limitations of requiring longer texts and greater computational effort, NERO yields a model-agnostic statistic that enables best-in-class detection, and principled, learning-free tracking of generative behavior over time.

Materials and Methods

Data acquisition. Human-authored texts were obtained from Project Gutenberg (English-language compositions) and technical manuscripts from arXiv via the official OAI-PMH interface. AI-generated texts were produced via webform and API access to the LLMs listed in Table 1. Details on prompt design, API access, and NERO computation are provided in the Supplementary Methods.

Code availability. A reproducible implementation of the entropy-rate estimator and end-to-end evaluation pipeline is available under a non-commercial license at <https://github.com/zeronowledgediscovery/nero>, and at zenodo (<https://doi.org/10.5281/zenodo.18487995>).

ACKNOWLEDGMENTS. This work was supported in part by the Defense Advanced Research Projects Agency (DARPA) under the ARC program (HR0011-26-3-E016). Views, opinions, and findings expressed are solely those of the authors.

1. Prepare for truly useful large language models. *Nat. Biomed. Eng.* **7**, 85–86 (2023).
2. E Crothers, N Japkowicz, HL Viktor, Machine-generated text: A comprehensive survey of threat models and detection methods. *IEEE Access* (2023).
3. AJ Thirunavukarasu, et al., Large language models in medicine. *Nat. medicine* **29**, 1930–1940 (2023).
4. TM Cover, RC King, A convergent gambling estimate of the entropy of english. *IEEE Transactions on Inf. Theory* (1978).
5. E Mitchell, Y Lee, A Khazatsky, CD Manning, C Finn, Detectgpt: Zero-shot machine-generated text detection using probability curvature in *Proceedings of the 40th International Conference on Machine Learning (ICML)*, Proceedings of Machine Learning Research, Vol. 202, pp. 24950–24962 (2023).
6. J Su, TY Zhuo, D Wang, P Nakov, Detectilm: Leveraging log rank information for zero-shot detection of machine-generated text in *Findings of the Association for Computational Linguistics: EMNLP 2023*. (2023).
7. P Gacs, J Tromp, PMB Vitányi, Algorithmic statistics. *IEEE Trans. Inf. Theory* **47**, 2443–2463 (2001).
8. P Grassberger, Estimating the information content of symbol sequences and efficient codes. *IEEE Transactions on Inf. Theory* **35**, 669–675 (1989).
9. I Chattopadhyay, H Lipson, Computing Entropy Rate Of Symbol Sources & A Distribution-free Limit Theorem. *ArXiv e-prints* 1401.0711 (2014).
10. I Chattopadhyay, H Lipson, Abductive learning of quantized stochastic processes with probabilistic finite automata. *Philos Trans A* **371**, 20110543 (2013).
11. CK Williams, CE Rasmussen, *Gaussian processes for machine learning*. (MIT press Cambridge, MA) Vol. 2, (2006).
12. OpenAI, Gpt-4 technical report (2023).
13. LA Levin, Laws of information conservation (non-growth) and aspects of the foundation of probability theory. *Probl. Inf. Transm.* **10**, 206–210 (1974) Originally in Russian.
14. M Li, PMB Vitányi, *An Introduction to Kolmogorov Complexity and Its Applications*. (Springer, New York, NY), 3rd edition, (2008).

373	15. GJ Chaitin, A theory of program size formally identical to information theory. <i>J. ACM</i> 22 , 329–340 (1975).	435
374	16. Y Horibe, A note on kolmogorov complexity and entropy. <i>Appl. Math. Lett.</i> 16 , 1129 – 1130 (2003).	436
375		437
376		438
377		439
378		440
379		441
380		442
381		443
382		444
383		445
384		446
385		447
386		448
387		449
388		450
389		451
390		452
391		453
392		454
393		455
394		456
395		457
396		458
397		459
398		460
399		461
400		462
401		463
402		464
403		465
404		466
405		467
406		468
407		469
408		470
409		471
410		472
411		473
412		474
413		475
414		476
415		477
416		478
417		479
418		480
419		481
420		482
421		483
422		484
423		485
424		486
425		487
426		488
427		489
428		490
429		491
430		492
431		493
432		494
433		495
434		496

Supplementary Methods

Proof of Proposition 1.

Proof. For typical $x \sim G_i$ (7, 14),

$$K(x) = K(S_i) + \log |S_i| + O(1). \quad [6]$$

Taking expectations over $x \sim G_i$, we have

$$\mathbb{E}_{x \sim G_i}[K(x)] = K(S_i) + \log |S_i| + O(1). \quad [7]$$

Subtracting, we find

$$\begin{aligned} & \mathbb{E}_{x \sim G_1}[K(x)] - \mathbb{E}_{x \sim G_2}[K(x)] \\ &= K(S_1) - K(S_2) + \log |S_1| - \log |S_2| + O(1). \end{aligned} \quad [8]$$

Invoking our assumption (Eq. Eq. (2)), this simplifies to

$$\mathbb{E}_{x \sim G_1}[K(x)] - \mathbb{E}_{x \sim G_2}[K(x)] = K(S_1) - K(S_2) + O(1).$$

This completes the first part. Noting that expected Kolmogorov complexity of strings from a stationary source satisfies (14, 16):

$$\lim_{n \rightarrow \infty} \frac{1}{n} \mathbb{E}_{x \sim G_i}[K(x)] = H(G_i). \quad [9]$$

establishes the second part. \square

Nonparametric entropy-rate estimation. We estimate the entropy rate of a stationary ergodic symbol stream using the PFSA-based nonparametric estimator introduced by Chattopadhyay *et al.* (9, 10). Given a sequence $s = s_1 \cdots s_n$ over a finite alphabet Σ , the method constructs empirical next-symbol distributions conditioned on observed substrings and identifies states of an a priori unknown underlying PFSA generator as equivalence classes of histories that lead to similar future distributions. Rare substrings are excluded via a frequency threshold m , ensuring statistical reliability without assuming a parametric model or access to the underlying generator. The entropy rate is then estimated as a weighted average of the Shannon entropies of empirical symbol-generation distributions associated with that context, with the weights inferred as observed frequencies of the classes of histories that lead to individual states. The estimator is provably consistent under stationarity and ergodicity, requires no training or reference model, and operates directly on the observed text stream. Full algorithmic details and theoretical guarantees are provided in (9, 10).

Pseudocode and Runtime Complexity. To enhance robustness and mitigate sensitivity to internal thresholds, we apply this estimator M times across a range of substring frequency thresholds, and report the median of the resulting entropy estimates, improving concentration guarantees (Theorem 1), with the error probability decaying exponentially in M .

Runtime complexity, including threshold-based pruning and median aggregation over M thresholds, is $O(nM|\Sigma|)$, where n is the length of the observed stream and $|\Sigma|$ the alphabet size (9). Thus, for a fixed M , the estimate has input-linear runtime complexity.

Algorithm 1 Robust Entropy-Rate Estimation

Input: Symbol stream x_1^n , alphabet Σ , threshold range $\{m_1, m_2, \dots, m_M\}$
Output: Robust entropy-rate estimate \hat{H}_n
foreach $m \in \{m_1, m_2, \dots, m_M\}$ **do**
 | Compute $\hat{H}_n^{(m)}$ with Chattopadhyay (9) with threshold m (ignore substrings with $< m$ occurrences)
 | Compute $\hat{H}_n \leftarrow \text{median}(\hat{H}_n^{(m_1)}, \hat{H}_n^{(m_2)}, \dots, \hat{H}_n^{(m_M)})$
return \hat{H}_n

Theorem 1 (Robustness and Concentration of the Aggregated Entropy Estimate). *Let x_1^N be a finite realization from a stationary, ergodic stochastic process over a finite alphabet Σ . Let $\hat{H}_n^{(m)}$ denote the entropy-rate estimate obtained by applying the algorithm of Chattopadhyay *et al.* (9, 10) with threshold m (substrings occurring $< m$ times ignored). Define the aggregated estimator*

$$\hat{H}_n := \text{median}\left\{\hat{H}_n^{(m_1)}, \hat{H}_n^{(m_2)}, \dots, \hat{H}_n^{(m_M)}\right\}.$$

If the indicators Z_i defined below are independent, then

$$\mathbb{P}(|\hat{H}_n - H_\star| > \epsilon) \leq \exp\left(-2M\left(\frac{1}{2} - \delta\right)^2\right).$$

More generally, when the $\hat{H}_n^{(m_i)}$ are obtained from randomly drawn length- L substreams of a fixed length- N text and hence may be dependent due to overlap, the same argument applies with an effective number of approximately independent draws

$$M_{\text{eff}} \asymp \frac{N}{L},$$

so that

$$\mathbb{P}(|\hat{H}_n - H_\star| > \epsilon) \lesssim \exp\left(-2M_{\text{eff}}\left(\frac{1}{2} - \delta\right)^2\right).$$

Proof. It follows from Chattopadhyay *et al.* (9, 10) that for sufficiently large n , for each threshold $m \in \{m_1, m_2, \dots, m_M\}$,

$$\mathbb{P}\left(|\hat{H}_n^{(m)} - H_\star| > \epsilon\right) \leq \delta,$$

for some fixed $\epsilon > 0$ and $\delta < \frac{1}{2}$. For each m_i , define

$$Z_i = \mathbf{1}\left\{|\hat{H}_n^{(m_i)} - H_\star| > \epsilon\right\}, \quad \mathbb{E}[Z_i] \leq \delta.$$

The median estimator \hat{H}_n deviates by more than ϵ only if at least half of the individual estimates do:

$$S = \sum_{i=1}^M Z_i \geq \frac{M}{2}.$$

If Z_1, \dots, Z_M are independent, Hoeffding's inequality yields

$$\mathbb{P}\left(S \geq \frac{M}{2}\right) \leq \exp\left(-2M\left(\frac{1}{2} - \delta\right)^2\right),$$

which implies the stated bound.

If instead the estimates are computed from randomly drawn length- L substreams of a fixed length- N text, the resulting Z_i may be dependent due to overlap. In that case, the same calculation can be interpreted in terms of an effective number of approximately independent draws $M_{\text{eff}} \asymp N/L$ (i.e., replacing M by M_{eff}), yielding the stated overlap-limited bound. \square

Classifier design (trained-NERO). For each document we form a fixed-dimensional feature vector $x \in \mathbb{R}^M$ from the PFSA entropy-rate estimator by evaluating the estimate at a prescribed set of substring-frequency cutoffs $\{m_1, \dots, m_M\}$. Specifically, the j th feature is the entropy-rate estimate computed using cutoff m_j , so that $x_j = \hat{H}(m_j)$. Missing feature values (arising when insufficient support is available at a given cutoff) are imputed with zeros using a constant-value imputer.

We evaluate several standard classifiers on these feature vectors, including random forests, extremely randomized trees, AdaBoost, gradient boosting, support vector machines (SVM), and Gaussian-process (GP) classification. Unless otherwise noted, models are fit on a random 50/50 train-test split of the pooled corpus, and performance is quantified by ROC/AUC on the held-out split. For tree-based models we use class-balanced weighting and 1000–5000 component estimators to stabilize performance; for SVM we use probabilistic calibration to obtain class probabilities.

We use GP classification as the trained-NERO readout as it yields the best performance under cross-validation. The GP classifier is implemented with an RBF kernel, $k(x, x') = \sigma^2 \exp(-\|x - x'\|^2/(2\ell^2))$, using the standard `GaussianProcessClassifier` implementation (11).

Birthtime trajectory fitting (Fig. 1c). To summarize how estimated entropy rate of generated text changes across successive LLM models, we fit a parametric growth curve to the model-specific time series shown in Fig. 1c. Each model is represented by a timestamp (“birthtime”) corresponding to its release date, the mean NERO estimate over generated tests, and an estimated standard deviation (error bar). Birthtimes were converted to elapsed time in days relative to the earliest cohort, $x_i = (t_i - t_0)$, yielding observations (x_i, y_i, σ_i) .

We fit a Richards (generalized logistic) curve with a fixed asymptotic upper bound U (corresponding to mean NERO estimate for human prose), parameterized as

$$f(x; L_0, k, x_0, \nu) = L_0 + \frac{U - L_0}{(1 + \exp\{-k(x - x_0)\})^{1/\nu}},$$

where L_0 is the lower asymptote, k the growth rate, x_0 the inflection location, and ν the shape parameter.

Parameters $\theta = (L_0, k, x_0, \nu)$ were estimated by weighted nonlinear least squares using `scipy.optimize.curve_fit` with heteroscedastic weights given by the cohort standard deviations σ_i . We used the initialization $(L_0, k, x_0, \nu) = (\max\{0, \min_i y_i - 0.03\}, 0.006, \text{median}(x), 0.6)$ and box constraints $L_0 \in [0, \min_i y_i]$, $k \in [10^{-6}, 3]$, $x_0 \in [-2000, 5000]$, $\nu \in [0.05, 5]$. The fitted parameters and their asymptotic standard errors were obtained from the returned covariance estimate $\widehat{\text{Cov}}(\hat{\theta})$, and standard chi-square diagnostics were computed from weighted residuals.

Data processing. For each corpus (human-authored and AI-generated), a structured CSV file was constructed at the document level. Each row corresponds to a single novel and contains the following fields: the entropy-rate estimate computed by NERO, the perplexity-based detector score reported by HowGPT, and the estimated percentage of AI-generated text reported by ZeroGPT. These tables form the basis for all quantitative comparisons reported in the main text and figures.

Topics and text generation. Topics used to prompt the AI models were obtained from a publicly available repository of narrative prompts (<https://www.plot-generator.org.uk/story-ideas/>). The precise semantic content of individual prompts is not critical for long-form text generation; rather, it is sufficient that prompts span a broad range of genres to avoid systematic topical bias. All topics were stored in a JSON file and loaded programmatically during text generation.

Each AI model was prompted with the identical ordered list of topics. This design ensures that differences in estimated entropy rate across models are not attributable to differences in prompt content.

AI text generation prompts. All AI-generated novels were produced using a fixed prompting protocol designed to elicit long-form narrative prose while minimizing structural or stylistic constraints beyond those necessary for length control and coherence. Identical prompt templates were used across models. Text generation was performed either via official REST APIs or, in some cases, via the OpenAI web interface using the same prompts and default generation settings; no prompt content differed between access modes.

Initial generation prompt. Each novel was initiated using the following base prompt, with bracketed fields instantiated programmatically:

I want you to act as a novelist writing about {topic}. The total length of the novel is about {novel_desired_length} characters. After generating each section of the novel, I will tell you how much text you have generated so far. In the novel, include plot development, characters, dialogue, and an overall narrative arc. Generate only the text of the novel itself, without annotations, explanations, or commentary. Write in full paragraphs and chapters of approximately {chapter.length} characters. Continue generating Chapter 1 until I tell you to stop. Generate exclusively novel text, with no metadata, confirmations, or chapter labels.

Continuation prompt. To extend generation beyond the initial context window and reach the target length, generation proceeded iteratively using a continuation prompt. At each step, previously generated text (truncated as necessary) was prepended, followed by the continuation instruction:

{previously generated text}

Here is the text generated so far. You have generated approximately {novel.current_length} characters. Continue writing Chapter {current.chapter} of the novel. Maintain narrative continuity and write in complete paragraphs. Do not generate annotations, explanations, summaries, or metadata. Generate exclusively novel text.

AI text generation implementation details. AI-generated novels were produced using an automated, iterative prompting pipeline implemented in Python. Each novel was generated toward a target length of 150,000 characters using repeated completion calls under default model parameters; no temperature, nucleus sampling, or penalty parameters were explicitly set. Only the maximum completion length was controlled.

Chapter indices were inferred implicitly as $|L/\ell| + 1$, where L is the cumulative character length generated so far and $\ell = 15,000$ is the nominal chapter length. Chapters were not explicitly labeled in the generated text.

To support long-form generation within finite model context windows, a sliding-window strategy was employed. When the concatenated prompt exceeded the available context length, only the most recent suffix of the previously generated text was retained, ensuring that the prompt plus maximum completion length remained within the model’s context window.

Generation continued until the target length was reached or until the model declined to continue. Responses yielding fewer than 200 output tokens were treated as refusals; in such cases, all text associated with that generation attempt was discarded and excluded from analysis.

API usage was subject to rate limits on requests per minute and token throughput. When limits were exceeded, generation was paused briefly before resuming. For each novel, all raw API responses and the cumulative generated text were stored in per-document metadata files to enable reproducibility and auditability. The generation pipelines are implemented in the following notebooks:

1. `gemini`: https://github.com/zeroknowledgediscovery/nero/api_data.collection/gemini-2.5.pro/get_novels.ipynb
2. `gpt4o`: https://github.com/zeroknowledgediscovery/nero/api_data.collection/openai/gpt4o/get_novels.ipynb
3. `gpt5`: https://github.com/zeroknowledgediscovery/nero/api_data.collection/openai/gpt5/get_novels.ipynb
4. `claude`: https://github.com/zeroknowledgediscovery/nero/api_data.collection/claude.sonnet.4/get_novels.ipynb

LLM access. AI-generated novels were produced using REST API access or web-interface access to the following models and versions:

- **Gemini**: `gemini-2.5-pro`, accessed between October 8, 2025 and October 16, 2025.
- **Claude**: `claude-sonnet-4-20250514`, accessed between August 26, 2025 and September 22, 2025.
- **GPT-4o**: `gpt-4o-2024-08-06`, accessed between November 13, 2025 and November 14, 2025.
- **GPT-5**: `gpt-5-2025-08-07`, accessed between October 26, 2025 and November 8, 2025.

API versus web-interface access. To mitigate potential bias arising from access modality, AI-generated texts were obtained using a mix of official REST API access and provider web interfaces, as summarized in Table 1. This design ensures that the reported entropy-rate differences are not artifacts of a particular deployment channel, prompting wrapper, or hidden system configuration. In particular, GPT-4.0 texts were generated using both API-based access and the OpenAI web interface under identical prompting protocols. We observed no statistically significant difference in estimated entropy rates between GPT-4.0 outputs obtained via API versus web-interface access.

Human-authored text corpus. Human-authored novels were retrieved from the Project Gutenberg corpus (<https://gutenberg.org/cache/epub/feeds/>), which provides full-text documents together with structured metadata. All available text and metadata files were downloaded, and only English-language works were retained based on metadata

745 fields. Legal headers and boilerplate text were removed using the
746 `gutenbergpy` library. Novels shorter than 150,000 characters after
747 preprocessing were excluded to ensure sufficient length for stable
748 entropy-rate estimation. 807
749 Technical manuscripts were collected from arXiv via the official
750 OAI-PMH interface and converted from `.tex` source to plain text.
751 These texts were used as a distinct human-authored reference corpus.
752 808
753 **Baseline detectors.** The HowGPT (<https://howkgpt.nyuad.nyu.edu/>) and
754 the ZeroGPT detectors (<https://www.zerogpt.com/>) were accessed between
755 November 2, 2025 and November 25, 2025, with access dates varying by
756 model cohort. From these tools, the reported perplexity-based detector
757 score (HowGPT) and the estimated percentage of AI-generated text
758 (ZeroGPT) were recorded. 809
759 For baseline detector comparisons, the central 15,000 characters of
760 each novel were extracted and submitted to two deployed detection
761 tools: HowGPT and ZeroGPT. The central segment was used to avoid
762 residual legal headers in human-authored texts; the same procedure was
763 applied to AI-generated texts to maintain parity. HowGPT/ZeroGPT
764 were queried on fixed-length segments due to interface constraints.
765 NERO was evaluated on full texts to reflect the intended long-form
766 detection regime. 810
767 **Perplexity baseline (HowGPT).** The perplexity-based baseline was
768 implemented using the HowGPT framework, which computes token-
769 level likelihood statistics under a reference language model and
770 aggregates them into a scalar detector score. We use the reported
771 detector score directly rather than raw perplexity values. 811
772 **ZeroGPT baseline.** The ZeroGPT baseline was obtained using the
773 deployed ZeroGPT web interface as described above, which returns
774 an estimated fraction of the submitted text labeled as AI-generated
775 (reported as a percentage). We record the percentage score exactly
776 as returned by the interface for each submitted segment, without
777 additional calibration or post-processing. This returned score was used
778 to construct the ROC curve. Because ZeroGPT is a proprietary system
779 whose underlying model, decision rule, and versioning may change over
780 time, the score should be interpreted as an external detector output at
781 the time of access rather than as a reproducible intrinsic statistic of
782 the text. 812
783 **Relation to standard LLM benchmarks.** Existing LLM benchmark
784 suites primarily evaluate task-specific performance (e.g., reasoning
785 accuracy, problem solving, or code synthesis) under controlled prompts
786 and scoring rules. In contrast, NERO estimates a distributional
787 property of long-form generated text—entropy rate under a fixed
788 symbolization—which reflects output complexity rather than task
789 competence. These quantities are not directly commensurate, and
790 strong correspondence is neither assumed nor required: models may
791 score highly on benchmarks while exhibiting low entropy-rate regularity
792 in extended prose. Accordingly, we do not attempt exhaustive
793 benchmark comparisons here. Studying how entropy-rate trajectories
794 relate to aggregate benchmark indicators across model releases is an
795 interesting direction for future work. 813
796
797
798
799
800
801
802
803
804
805
806 814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868